

# Decoded Quantum Interferometry: Interference-Driven Quantum Speedup

*Optimization as Decoding for Fourier-Sparse Objectives*



**PHYS360 — Paper Draft**

Instructor: Dr. Chandra Turpen

Student: Jaehyun Kwak  
UID: 122274570 · Email: [apjh0703@terpmail.umd.edu](mailto:apjh0703@terpmail.umd.edu)

Affiliation: University of Maryland, College Park (Exchange from Korea University)

Word Count (text-only, excluding references/figures): \_\_\_\_\_  
Draft/Final: **Draft**

# Table of Contents

- I. Introduction
- II. Background
  - A. Quantum Computing Basics
    - i. Superposition and Interference
    - ii. Interference for Global Structure Extraction
  - B. Prior Work
    - i. Hamiltonian-Based Approaches
    - ii. Limitations of Hamiltonian-based Approaches
- III. Decoded Quantum Interferometry
  - A. Reformulating Optimization as Decoding
  - B. Fourier Sparsity and Structural Promise
  - C. DQI Analysis
- IV. Significance and Potential Impact
  - A. Broader Class of Applicable Problems
  - B. Potential Practical Relevance
- V. Limitations
  - A. Hardware Constraints
  - B. Possibility of Classical Catch-Up
- VI. Conclusion
- References

## I. Introduction

This review paper introduces a recent paper that proposes **Decoded Quantum Interferometry (DQI)** as an alternative way to tackle discrete optimization.

Among those quantum optimization methods that were suggested to overcome their classical counterparts, DQI is distinctive as it diverges from conventional methodologies.

Optimization shows up in almost every field of science and engineering. From routing to model tuning, whenever there is a need to choose the best action under constraints, optimization offers a strong strategy for a solution. However, when it comes to scalability, it sacrifices accuracy, as the search space typically grows combinatorially. Although some classical methods such as convex relaxation and heuristics have been devised to cope with this limit, their performance often deteriorates as dimensionality, noise, or hard combinatorial interactions increase.

Quantum computing suggests an alternative approach. In terms of complexity, BQP (Bounded-error Quantum Polynomial time) captures problems solvable efficiently by quantum algorithms with small error. Yet, quantum advantage does not imply an unconditional speed up but appears when the problem structure aligns with quantum resources. One of the representatives that satisfies this condition is Shor's factoring algorithm [1], which also motivated DQI.

In the realm of optimization, most prior quantum approaches are Hamiltonian-based [2], in which they encode a cost function as an energy landscape and attempt to reach low-energy states via quantum annealing or the Quantum Approximate Optimization Algorithm (QAOA). They are intuitive and broadly applicable, but at the same time they have practical limitations with circuit depth, calibration noise, barren plateaus, and the difficulty of parameter training on near-term devices.

In this paper a newer perspective is introduced – Decoded Quantum Interferometry (DQI). Instead of exploiting Hamiltonian-based, energy minimization, DQI converts optimization problems into decoding latent structure. The circuit uses interference to imprint global patterns of the objective into measurement statistics. Then a lightweight classical routine decodes those statistics to recover informative features.

This paper mainly aims for explaining DQI and giving a simple analysis of it. We begin with reviewing some fundamental backgrounds on quantum computing – superposition and interference. Then we summarize prior works that mainly focused on Hamiltonian-based approaches and show their limitations. Next, we analyze DQI and discuss its potential applications and limitations it faces.

## II. Background(864)

Before deep diving into DQI and its analysis, for better understanding it is desirable to review the most fundamental quantum computing basics and how prior studies approached optimization problem in the field of quantum computing. This section briefly revisits superposition and interference and then introduces prior approaches using Hamiltonian-based algorithms with their limitations.

### A. Quantum Computing Basics

Among distinctive features of quantum computers, superposition and interference are the two most notable ones, as they form the basis of how quantum algorithms encode and reveal structure.

#### i. Superposition and Interference

Superposition is a phenomenon that stems from quantum mechanics. It enables a quantum system to be in multiple states simultaneously. It lets a single qubit occupy a complex-weighted blend of basis states, typically written  $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$  with  $|\alpha|^2 + |\beta|^2 = 1$ . Here the weights, or amplitudes, determine measurement probabilities via their norm-squares. Measuring a qubit, an outcome becomes definite among the states – “Measurement collapses the qubit.” [3, p. 85]

Interference is a quantum mechanical phenomenon that emerges for electrons and photons having wavelike behavior. [4] When a single particle is in a superposition, its amplitudes can add or subtract, yielding constructive or destructive interference. Relative phase is the key, as it redirects probability weight, changing how components overlap with the measurement basis. On the other hand, as probabilities come from squared magnitudes, a global phase that multiplies every amplitude equally is physically irrelevant.

#### ii. Interference for Global Structure Extraction

Outcomes strengthen with constructive interference and weaken when destructive, so by reading the interference pattern we can infer the problem’s global structure. A set of simple examples can be delivered by Hadamard gate. Given  $|0\rangle$  an initial state, by applying Hadamard gate we get a state in superposition as followed.  $H|0\rangle = (|0\rangle + |1\rangle)/\sqrt{2}$ . Measuring with Z-basis, this state may collapse into either  $|0\rangle$  or  $|1\rangle$ , with equal probability of  $1/2$ . Applying Hadamard gate once more,  $H(H|0\rangle) = (1/\sqrt{2})(H|0\rangle + H|1\rangle) = (1/2)((|0\rangle + |1\rangle) + (|0\rangle - |1\rangle)) = |0\rangle$ , the

state cancels out the probability of becoming  $|1\rangle$ , and ends up being  $|0\rangle$ , strengthening its probability. On the other hand, by simply adding Z gate in the middle we can discover  $|0\rangle$  being cancelled out, by destructive interference.  $HZH|0\rangle = HZ(|0\rangle + |1\rangle)/\sqrt{2} = H(|0\rangle - |1\rangle)/\sqrt{2} = (1/\sqrt{2})(H|0\rangle - H|1\rangle) = (1/2)((|0\rangle + |1\rangle) - (|0\rangle - |1\rangle)) = |1\rangle$ . Here the  $|0\rangle$  is cancelled out, the state resulting in  $|1\rangle$ .

One notable application utilizing this is Grover's search. Grover's search is a searching algorithm that exploits constructive interference. Suppose there is one marked item among  $N$  basis states, whose phase is flipped by an oracle, while others unchanged. The diffusion operator then inverts amplitudes about their mean, recombining paths so the probability of the marked state strengthens and the rest diminishes. Each diffusion round by the oracle performs a rotation toward the marked basis vector by roughly  $2/\sqrt{N}$ . Approximately after  $(\pi/4)\sqrt{N}$  rounds, the probability of the marked state becomes high enough, giving  $O(\sqrt{N})$  queries instead of classical  $O(N)$ . [5] Exploiting-large scale constructive interference, Grover's search identifies the global structure and yields a speedup.

## B. Prior Work

Before DQI, there have been many attempts to develop optimization algorithms using quantum computing. In this section, those efforts, their limitations and what makes DQI different are introduced.

### i. Hamiltonian-Based Approaches

Hamiltonian-based approach has been the dominant line of work in the field of quantum optimization. It encodes a combinatorial objective as a **Hamiltonian**, whose low-energy states correspond to ideal solutions. There are two main ways of implementing – adiabatic/annealing methods and QAOA, which refers to Quantum Approximate Optimization Algorithm. An adiabatic/annealing method uses continuous time. With a mixer  $H_m$  that spreads amplitude and induces transitions, it prepares the ground state of the mixer and slowly interpolate  $H(s) = (1 - s)H_m + sH_c$ . When the evolution is slow relative to the minimum spectral gap, the system is near the ground state of  $H(s)$  and ends close to the optimum of  $H_c$ , which is an ideal solution.

On the other hand is QAOA, a variational quantum algorithm for combinatorial optimization. It is a hybrid method combining short quantum circuits and a classical optimizer. QAOA consists of mainly two

steps and repeats them a few times. First it turns the problem into a simple score function and repeats two steps a few times: a cost step that tags candidate answers with phases based on their score, and a mixer step that shuffles amplitude so these phases interfere. After each round it measures the qubits, estimates the average score, and let a classical optimizer predicts the optimal parameters – angles of qubits – for the next round.

These methods are attractive in various aspects. The “cost-as-energy” picture is intuitive, mapping methods are mature, and most importantly they are hardware-friendly for their shallow-depth variants, for small instances.

## ii. **Limitations of Hamiltonian-Based Approaches**

However, there are practical limitations of this methodology. To improve quality, adiabatic models or QAOA usually need longer schedules or deeper circuits, which bring about problems of coherence and gate-noise. This is critical in that even the latest quantum computers struggle with stability and scale. Moreover, perfect hardware by itself would not guarantee feasibility as problems with small spectral gaps - energy gap above the ground state - force very slow adiabatic evolution. Remedies such as minor embedding also adds qubits and noise. Variational QAOA loops require many samples to estimate the cost and can stall during training due to barren-plateaus. Calibration drift also changes device behavior over time, so parameters that worked before may stop working.

Overall, performance is highly instance-dependent and there rarely is a strong guarantee. This is where DQI jumps in. Avoiding deep energy minimization, it uses short interference blocks and a lightweight classical decoder to tackle the problem.

## **III. Decoded Quantum Interferometry**

In contrast to Hamiltonian methods that focus on local cost landscapes, DQI operates on the objective’s Fourier representation, letting it exploit spectral structure more if present. Fourier representation writes a function as a sum of simple building blocks: frequencies. For Boolean functions the blocks are parity characters, whose weights are called Fourier coefficients. Bigger coefficients imply stronger interactions among the corresponding variable sets. DQI’s interference patterns make those terms stand out in the statistics when measured.

Here is a rough picture how DQI works. First it prepares a shallow superposition, applies a few interference blocks that phase-marks basis states according to the

objectives, measuring repeatedly. Then the patterns reveal which parity features matter. A lightweight classical decoder is used to extract the dominant coefficients, that are later used for sampling good solutions.

In this section we first focus on the concept of Fourier sparsity and its structural promise, as they are the foundational block of DQI. Then we discover how DQI reformulates optimization problems as decoding. Lastly, referring to the original paper, the whole picture of DQI is given.

## A. Fourier Sparsity and Structural Promise

Many combinatorial objectives have a simple shape when viewed in the Fourier basis. Instead of thinking in terms of raw variables, this view focuses on parity features, which are products of a few variables. In many problems, only a small number of such features matter, and usually the important ones involve only few variables. In other words the spectrum is sparse and low-degree, and that aligns well with DQI's Fourier-based approach – though not universally guaranteed. Here is the basic overview of Fourier representation.

For a Boolean objective  $f: \{\pm 1\}^n \rightarrow R$ , we can rewrite  $f$  as a sum of simple parity features:

$$f(x) = \sum_{S \subseteq [n]} \hat{f}(S) \chi_S(x), \quad \chi_S(x) = \prod_{i \in S} x_i.$$

Each  $\chi_S$  multiplies a small set of variables, and the magnitude of  $S$  refers to the number of variables under effect. The number  $\hat{f}(S)$  is the Fourier coefficient, which is the weight of that feature. A sparse and low-degree spectrum means that only a few coefficients  $\hat{f}(S)$  are large, mostly with small  $|S|$ .

One good practical application of Fourier representation is the Max-cut problem. [2] The Max-cut problem asks one to divide a graph's vertices into two groups, so that as many edges as possible run between them – in other words, as many edges as possible are cut by the partition boundary.

Let  $G = (V, E)$  with  $|E| = m$ . Give each vertex a label  $x_i \in \{\pm 1\}$ , which indicates which side of the cut it sits on. An edge  $(i, j)$  is cut when they are in different groups, i.e.,  $x_i x_j = -1$ . When unweighted, Max-cut objective function is written as

$$f(x) = \frac{1}{2} \sum_{(i,j) \in E} (1 - x_i x_j) = \frac{m}{2} - \frac{1}{2} \sum_{(i,j) \in E} x_i x_j, \quad m = |E|$$

In the  $\pm 1$  encoding,  $x_i x_j = +1$  when vertices  $i$  and  $j$  are on the same side and  $-1$  when they are separated. Thus  $(1 - x_i x_j)$  equals 0 for an uncut edge and 2 for a cut edge, so  $\frac{1}{2}(1 - x_i x_j)$  is exactly the indicator that edge  $(i, j)$  is uncut. Summing over all edges, we get the equation above. Now we compare it with the Fourier form  $f(x) = \sum_{S \subseteq [n]} \hat{f}(S) \chi_S(x)$ , where  $\chi_S(x) = \prod_{i \in S} x_i$ . The expansion above has a constant part  $\frac{m}{2}$ , giving  $\hat{f}(\emptyset) = \frac{m}{2}$ , for random partitioning an edge is cut with  $\frac{1}{2}$  probability. The other part is a sum of pairwise parities  $x_i x_j = \chi_{\{i, j\}}(x)$ , with coefficient  $\frac{1}{2}$  on each edge, so  $\hat{f}(\{i, j\}) = -\frac{1}{2}$ ,  $(i, j) \in E$ . Every other coefficient vanishes. There is no bias on individual variables and no higher-order terms, as the objective depends only on pair interactions along edges. Consequently, Max-cut's spectrum is exactly low-degree and sparse. The only two possible values for its degree are 0 and 2, as the degree means that only a small subset of coefficients carry significant weight, while most are exactly zero or negligibly small – indeed, in Max-Cut only the constant and one pairwise coefficient per edge are nonzero (support size  $1 + |E|$  out of  $2^n$ ). Taken together, Max-Cut provides a clear example of the low-degree, edge-aligned structure that DQI exploits.

## B. Reformulating Optimization as Decoding

DQI replaces Hamiltonian-based approaches' doctrine, “minimize an energy landscape” with “decode hidden structure from interference”. Given that  $f(x)$  is the objective function on bit strings, rather than searching for the lowest point, DQI asks which structure – e.g., parities / Fourier characters – determines good solutions, and how we can read them out by interference. Therefore, the circuit follows the logical structure: prepare  $\rightarrow$  interfere  $\rightarrow$  decode.

First, an  $x$ -register is initialized to be in the uniform superposition of  $|+\rangle^{(\otimes n)}$ . Then, an ancilla  $y$ -register that coherently ranges over small selections of low Hamming weight is prepared. Practically, this means that for each  $k \leq l$  we form an equal superposition of all basis states that contain exactly  $k$  ones – and  $n - k$  zeros, and then combine those fixed-weight sectors with chosen coefficients (a degree budget  $l$ ).

Second, to make them interfere phase marking is done, according to the objective so that components that agree with useful structure acquire aligned phase. In the  $\text{mod } p$  linear setting of DQI, the circuit computes a syndrome  $s = B^T y$  and imprints phases that depend on how  $x$  satisfies the selected constraints. Because amplitudes add with signs, a final layer of Hadamards on

the x-register converts these relative phases into amplitude differences, which is the interference pattern that encodes global information about  $f$ .

Third step is decoding. Using only the syndrome  $s$  and the known code  $B$ , a reversible wrapper of a classical decoder infers  $\hat{y}$ . XORing  $\hat{y}$  into the ancilla erases it:  $|y\rangle \rightarrow |y \oplus \hat{y}\rangle = |0 \dots 0\rangle$ . This leaves a clean state on the x-register whose amplitudes are a polynomial in the objective. Measuring  $|P(f)\rangle$  yields bit strings with probability proportional to  $P(f(x))^2$ . Thus, DQI biases across level sets of  $f$  while remaining unbiased within each level set. The degree  $P = l$  controls the amplification, decoding trade-off. Higher degree gives stronger bias but requires a decoder that corrects up to  $l$  errors.

## C. DQI Analysis

At heart, DQI turns optimization into amplitude shaping. The quantum circuit prepares a short superposition, uses interference to encode information about the objective into phases, and then a final Hadamard converts those phases into amplitudes so that the output state has the form  $|P(f)\rangle = \sum(x)P(f(x))|x\rangle$  [2]. Measuring this state produces samples with probability proportional to  $P(f(x))^2$ , so that choosing a higher-degree polynomial  $P$  strengthens the bias toward large objective values. But at the same time it also raises the required error-correction capability of the classical decoder to the same degree  $l = \deg P$  [2].

The original paper quantifies this bias for a broad family called max-LINSAT. If each constraint set has size  $r$  over  $F_p$  and the decoder can correct up to  $l$  bit-flip errors, the expected fraction of satisfied constraints after one DQI sample obeys the semicircle law,

$$\langle s \rangle / m = \left( \sqrt{\left( \frac{1}{m} * (1 - r/p) \right)} + \sqrt{\left( \frac{r}{p} * (1 - 1/m) \right)} \right)^2,$$

With the baseline  $r/p$  equal to the uniform-random satisfaction rate. At  $r/p = 1/2$  the expression is exactly a semicircle.

Two case studies illustrate how decoding power translates into optimization performance: Optimal Polynomial Intersection (OPI) and sparse max-XORSAT [2].

OPI is a univariate algebraic special case over a finite field. For each  $y$  we are given a small allowed set  $F_y$ , and we seek a low-degree polynomial  $Q$  that hits as many  $F_y$ 's as possible [2]. DQI maps this task to decoding a standard classical algebraic code, which  $x \in \{0,1\}^n$  must satisfy as many XOR equations

$\bigoplus_j B_{(i,j)} x_j = b_i$  as possible. DQI's quantum stage marks the superposition by how many parity checks each assignment satisfies and converts those markings into measurement patterns [2]. On the classical side, we use a standard, off-the-shelf iterative decoder for sparse parity check systems, which are the same style widely used in error-correction (LDPC), to infer a small set of it flips that would satisfy most checks [2]. When the constraint graph is genuinely sparse and well-conditioned, this decoder is fast and effective, so DQI can use a larger degree budget  $l$  and thus the amount of amplification [2].

Beyond approximation quality, DQI's sampling property matters, as it is fair within level sets. Among solutions with equal objective value, sampling is uniform. Unlike most classical heuristics that are path-dependent and introduce bias toward certain regions, DQI's fairness enables approximate counting by turning repeated samples into level-set multiplicity estimates.

## IV. Significance and Potential Impact

DQI excels when the objective exhibits low-degree, Fourier-sparse patterns that align with a good classical decoder. In such regime it delivers provable bias with semicircle law, competitive or superior approximation on OPI, and fair sampling for counting, as it makes clear how performance scales with the degree budget  $l$  and the strength of the chosen decoder.

### A. Broader Class of Applicable Problems

DQI is not tied to one constraint family because its quantum part only needs a syndrome and its classical part is a decoder. That means you can swap in decoders that match the mathematical structure of your constraints. DQI covers tasks like parity-based satisfiability like max-XORSAT, modular constraint satisfaction, sparse Ising/QUBO objectives including Max-Cut, and problems where constraints can be summarized by a short consistency check. The only two requirements are as followed: when one can compute a concise summary  $s = B^T y$  that reflects which constraints were touched and a fast classical routine – decoder - can recover a small selection  $y$  of involved constraints from that summary [2]. When those hold, DQI's interference highlights the right features, and the classical step cleans

up the rest. Weighted or filtered variants are handled by simple scaling/replication and pre/post-processing, without changing the quantum template.

## B. Potential Practical Relevance

On near-term devices, DQI's shallow prepare-interfere blocks reduce exposure to depth limits, calibration drift, and barren-plateau training, which have been main limitations for the Hamiltonian-based approaches. In addition, the output given by DQI is fair within level sets enabling approximate counting and unbiased benchmarking of solution diversity, which are properties uncommon in classical heuristics. In practice, DQI can serve as a biasing sampler that concentrates probability on high-quality selected assignments. Selected samples can then be refined by classical local search. The heavy part sits in mature classical decoders, making performance predictable, portable, and improvable, as decoding communities advance.

## V. Limitations

As mentioned, DQI full potential requires some preconditions to be met. The major limitations of this methodology are as follow: shot complexity to resolve statistics, device stability that guarantees reliable interference, and the availability of a decoder that matches the problem's structure. Additionally, as DQI's performance over classical counterparts are not mathematically strictly proven, there is still a chance for the classical community to follow it up.

## A. Hardware Constraints

Even though DQI uses shallow prepare-interfere blocks, it still needs stable, phase-coherent interference across many shots. Limited coherence time, calibration drift, crosstalk, and some other physical qubits' errors are a severe threat for DQI to operate properly. Shot complexity also matters. To resolve small differences between level sets, one may need to repeat sampling many times, which magnifies any slow drift. Finally, some decoders require modest ancilla width and reversible arithmetic. Today's hardwares are yet not mature enough to guarantee all those conditions.

## B. Possibility of Classical Catch-Up

DQI's advantage hinges on structure-plus-decoder fit. If the target lacks low-degree, sparse parity structure, or if a strong and fast classical routine also exploits it, classical methods can match or surpass DQI. Examples include tailored simulated annealing, as its score was even higher than that of DQI when run for long time, which briefly shows that DQI is not unconditionally superb. Moreover, pushing  $l$  higher strengthens bias but raises decoding difficulty. If decoding becomes the bottleneck, any end-to-end advantage shrinks, and some other strong competitors from both quantum and classical community may arise.

## VI. Conclusion

Decoded Quantum Interferometry recasts optimization as decoding: short prepare-interfere circuits shape amplitudes so measurements are biased by a polynomial of the objective yet remain fair within level sets. The analysis clarifies when structure and a matching decoder yield gains, with case studies from OPI to sparse XORSAT. Practically, DQI shifts depth and training burdens onto efficient classical routines, making it attractive for NISQ devices. Its limits including shot complexity, phase stability, and decoder availability, however suggest instance selection. Promising directions include denser constraints, stronger decoders, and quantum decoding, as well as hybrid pipelines that refine DQI samples with classical research. But it is clear that DQI has opened a new era in the field of quantum optimization. This paper would last as the textbook-like study for the following researchers.

## References

- [1] P. W. Shor, Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer, SIAM Journal on Computing, 1997.
- [2] S. P. Jordan, N. Shutty, M. Wootters, A. Zalcman, A. Schmidhuber, R. King, S. V. Isakov, T. Khattar and R. Babbush, Optimization by Decoded Quantum Interferometry, Nature, 2025.
- [3] T. G. Wong, Introduction to Classical and Quantum Computing, Omaha, Nebraska, 2022.
- [4] "Explore quantum Interference," Microsoft, [Online]. Available: <https://quantum.microsoft.com/en-us/insights/education/concepts/interference>. [Accessed 08 11 2025].
- [5] "Fundamentals of quantum algorithms," IBM Quantum Platform, [Online]. Available: <https://quantum.cloud.ibm.com/learning/en/courses/fundamentals-of-quantum-algorithms/grover-algorithm/introduction>. [Accessed 05 11 2025].